

Distributed Computing

MPI

AGENDA

PTP communication



MPICH

PTP examples

End
Questions




Knowing and installing MPICH

MPICH (www.mpich.org)

- MPICH is a **freely** available, **portable** implementation of **MPI**, a standard for message-passing for **distributed-memory** applications used in **parallel** computing.
- The CH part of the name was derived from "**Chameleon**", which was a portable parallel programming library developed by **William Gropp**, one of the **founders** of MPICH.
- History:
 - **Before 2001**: MPICH1 which implements MPI-1
 - **Between 2001-2012**: MPICH2 which implements MPI-2
 - **After 2012**: MPICH v3.0 which implements MPI-3

MPICH installation (For our lab)

- MPICH installation requirements:
 - Windows XP.
 - .NET framework 2.0.
 - VS2005.
- To install MPICH in your machine use the instructions that can be found in this link: <http://www.cs.utah.edu/~delisi/vsmpi/>
- After installation and compiling your MPI use this command on CMD:
 - `mpiexec -np 4 yourprogram.exe`



Point to point commuincation in
MPI

Communicators

- In MPI, all communication operations are executed using a **communicator**. A communicator represents a communication domain which is essentially a **set** of **processes** that exchange messages between each other.
- The MPI default communicator MPI COMM WORLD is used for the communication. This **communicator captures** all **processes** executing a **parallel** program

MPI_Send operation

- `int MPI_Send(void *smessage, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm) .`
- **smessage** specifies a send buffer which contain the data elements.
- **count** is the number of elements to be sent from the send buffer.
- **datatype** is the datatype of each entry in the send buffer.
- **dest** specifies the rank of the target process.
- **tag** is a message tag which can be used by the receiver to distinguish different messages from the same sender.

MPI_Recv operation

- `int MPI_Recv(void *rmessage,
 int count,
 MPI_Datatype datatype,
 int source,
 int tag,
 MPI_Comm comm,
 MPI_Status *status) .`
- Like MPI_Send except that:
 - **rmessage** specifies the receive buffer in which the message should be stored.
 - **status** specifies a data structure which contains information about a message after the completion of the receive operation. Can be ignored using `MPI_STATUS_IGNORE`

MPI_Recv operation Cont.

- Like MPI_Send except that:
 - By using **source** = MPI_ANY_SOURCE, a process can receive a message from any arbitrary process.
 - Similarly, by using **tag** = MPI_ANY_TAG, a process can receive a message with an arbitrary tag.
- After completion of the receive operation, status variable will contain these information:
 - **status.MPI_SOURCE** specifies the rank of the sending process.
 - **status.MPI_TAG** specifies the tag of the message received.
 - **status.MPI_ERROR** contains an error code.

MPI Datatypes

MPI Datentyp	C-Datentyp
MPI_CHAR	signed char
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_LONG_LONG_INT	long long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_UNSIGNED_LONG_LONG	unsigned long long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_WCHAR	wide char
MPI_PACKED	special data type for packing
MPI_BYTE	single byte value

What happens while sending and receiving ?

1. The data elements to be sent are **copied** from the send buffer **smessage** specified as parameter into a **system buffer** of the MPI runtime system. The **message** is **assembled** by adding a **header** with information on the **sending process**, the **receiving process**, the **tag**, and the **communicator** used.
2. The **message** is sent via the **network** from the **sending** process to the **receiving** process.
3. At the **receiving** side, the data entries of the **message** are **copied** from the **system buffer** into the **receive buffer** **rmessage** specified by `MPI_Recv()`.

Send and receive operations nature

- Both MPI_Send() and MPI_Recv() are blocking, asynchronous operations.
 - MPI_Recv() operation can also be **started** when the **corresponding** MPI_Send() operation has not yet been **started**.
 - The **process** executing the MPI_Recv() operation is **blocked** until the specified receive **buffer** contains the data elements sent.
 - Similarly, an MPI_Send() **operation** can also be **started** when the corresponding MPI_Recv() **operation** has not yet been **started**.



Point to point operation examples



QUESTIONS

THANK YOU!

